


# Structure learning of Bayesian networks with the PC (Peter–Clark) algorithm

*Modestas Filipavicius*

*May 3, 2018*

We're given a partially directed acyclic graph (pDAG), with 6 random variables:

```
partial_dag = new("graphNEL", nodes=c("A", "B", "C", "D", "E", "F"), edgemode="directed")
partial_dag = addEdge("A", "C", partial_dag, 1)
partial_dag = addEdge("C", "D", partial_dag, 1)
partial_dag = addEdge("B", "C", partial_dag, 1)
partial_dag = addEdge("B", "E", partial_dag, 1)
plot(partial_dag)
```




Within a Bayes Net with 6 random variables, are two variables, A and B, causally related?

We will run the PC algorithm to learn correct graph structure and causal relationships.

Plot the A and B observations:

```
data = readRDS("MVN_DAG.rds")
data$ID = seq.int(nrow(data))

library(ggplot2)
ggplot(data, aes(ID, y = value, color = variable)) +
  geom_point(aes(y = A, col = "A")) +
  geom_point(aes(y = B, col = "B"))
```



What does the plot suggest about their (marginal) correlation?

```
# calculate Pearson Correlation Coef (PCC) between A and B
cor_AB = cor.test(data[, 1], data[, 2])
cor_AB$p.value
```

```
## [1] 0.840103
```

```
# Null hypothesis, that A and B are not related, holds with high confidence.
```

```
# Just for fun, test for correlations between A and C, and B and C
cor_AC = cor.test(data[, 1], data[, 3])
cor_BC = cor.test(data[, 2], data[, 3])
cor_AC$p.value
```

```
## [1] 2.261867e-18
```

```
cor_BC$p.value
```

```
## [1] 6.901592e-32
```

```
# these p-values are extremely small, and we can reject
# null hypothesis with high confidence!
```

```
# "negative" control: unconnected node F to B
# expect very low correlation
```


```

cor_BF = cor.test(data[, 2], data[, 6])
cor_BF

##
## Pearson's product-moment correlation
##
## data: data[, 2] and data[, 6]
## t = -1.4007, df = 298, p-value = 0.1624
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.19235247 0.03266679
## sample estimates:
## cor
## -0.08087322

# calculate all pair-wise PCCs
cor_all = round(cor(data[, 1:5]), 2)
heatmap(cor_all, xlab = 'All pair-wise PCCs for 6 nodes')

```



All pair-wise PCCs for 6 nodes

## Testing for partial correlation

```

# Linearly regress A on C (that is, with A as the response variable
# and C as the explanatory variable).
# Compute and store the residuals.
lm_AC = lm(data[, 1] ~ data[, 3], data = data)

```


```

res_AC = residuals(lm_AC)

# Linearly regress B on C. Compute and store the residuals.
lm_BC = lm(data[, 2] ~ data[, 3], data = data)
res_BC = residuals(lm_BC)

# Plot the residuals of A (regressed on C) against
# the residuals of B (regressed on C). What do you see?
plot(x = res_AC, y = res_BC)

```



In the plot above, I can see a slight negative correlation between residuals test for Pearson correlation between A and B residuals regressed on C.

Let's calculate an exact value of Pearson Correlation:

```

cor_ABgC = cor.test(res_AC, res_BC)

# PCC is in fact negative:
cor_ABgC

##
## Pearson's product-moment correlation
##
## data: res_AC and res_BC
## t = -7.5173, df = 298, p-value = 6.6e-13
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.4903245 -0.2995546
## sample estimates:
## cor
## -0.3992521

# P-value
cor_ABgC$p.value

```

```
## [1] 6.599955e-13
```

The p-value is extremely small, thus we can reject the null hypothesis, namely, that there's no correlation between residuals, with high confidence.

## Running the PC algorithm to learn correct graph structure and causal relationships


```
# install.packages("pcalg")
# source("https://bioconductor.org/biocLite.R")
# biocLite("graph")
# biocLite("RBGL")
# biocLite("Rgraphviz")

library(pcalg)

# calculate covariance matrix
data_cor = cor(data[, 1:6])
data_n = nrow(data)
suffStat = list(C = data_cor, n = data_n)


# alpha: significance level (number in (0, 1) for the individual
# conditional independence tests.
a = pc(alpha = 0.10, suffStat = suffStat, indepTest = gaussCItest, labels = c(colnames(data[, 1:6])))
plot(a)
```

```
pc(suffStat = suffStat, indepTest = gaussCItest, alpha = 0.1,
  labels = c(colnames(data[, 1:6])))
```




```
a = pc(alpha = 0.20, suffStat = suffStat, indepTest = gaussCIttest, labels = c(colnames(data[, 1:6])))
plot(a)
```

**pc(suffStat = suffStat, indepTest = gaussCIttest, alpha = 0.2,  
labels = c(colnames(data[, 1:6])))**




```
a = pc(alpha = 0.40, suffStat = suffStat, indepTest = gaussCIttest, labels = c(colnames(data[, 1:6])))
plot(a)
```

```
pc(suffStat = suffStat, indepTest = gaussCIttest, alpha = 0.4,
  labels = c(colnames(data[, 1:6])))
```



```
a = pc(alpha = 0.9, suffStat = suffStat, indepTest = gaussCIttest, labels = c(colnames(data[, 1:6])))
plot(a)
```

```
pc(suffStat = suffStat, indepTest = gaussCIttest, alpha = 0.9,
  labels = c(colnames(data[, 1:6])))
```



We observe that the closest structure to figure 2 is obtained by setting  $\alpha = 0.1$ . Then the equivalence class joint distribution looks like  $p(a, b, c, d) = p(a)p(b)p(c|a, b)p(d|c)$  and is only missing E and F nodes, which are completely disconnected. Increasing  $\alpha$  to  $\alpha = 0.2$  connects E and F via undirected edge. As soon as  $\alpha = 0.25$ , every node is connected to at least one other node. However, I was not able to reproduce the figure 2 with any  $\alpha$  values.

Here  $\alpha$  param is defined as “significance level (number in  $(0, 1)$  for the individual conditional independence”, and roughly corresponds to reducing the dependency threshold needed to connect two graphs. In other words, it’s a regularization coefficient that we must optimize for a given dataset.